

Computational Photography

- * Study the basics of computation and its impact on the entire workflow of photography, from capturing, manipulating and collaborating on, and sharing photographs.



© 2015 Irfan Essa, Georgia Tech, All Rights Reserved

Making a Panorama



(Lords Cricket Ground, London, UK, by I. Essa)



Lesson Objectives

1. Generate a Panorama
2. Image Re-projection
3. Homography from a pair of images
4. Computing inliers and outliers
5. Details of constructing panoramas

Review: 5 Steps to Make a Panorama



(Lord's Cricket Ground, London, UK, by I. Essa)

- * Capture Images
- * Detection and matching
- * Warping → Aligning Images
- * Blending, Fading, Cutting
- * Cropping (Optional)

Align Images: Translate??



L on top

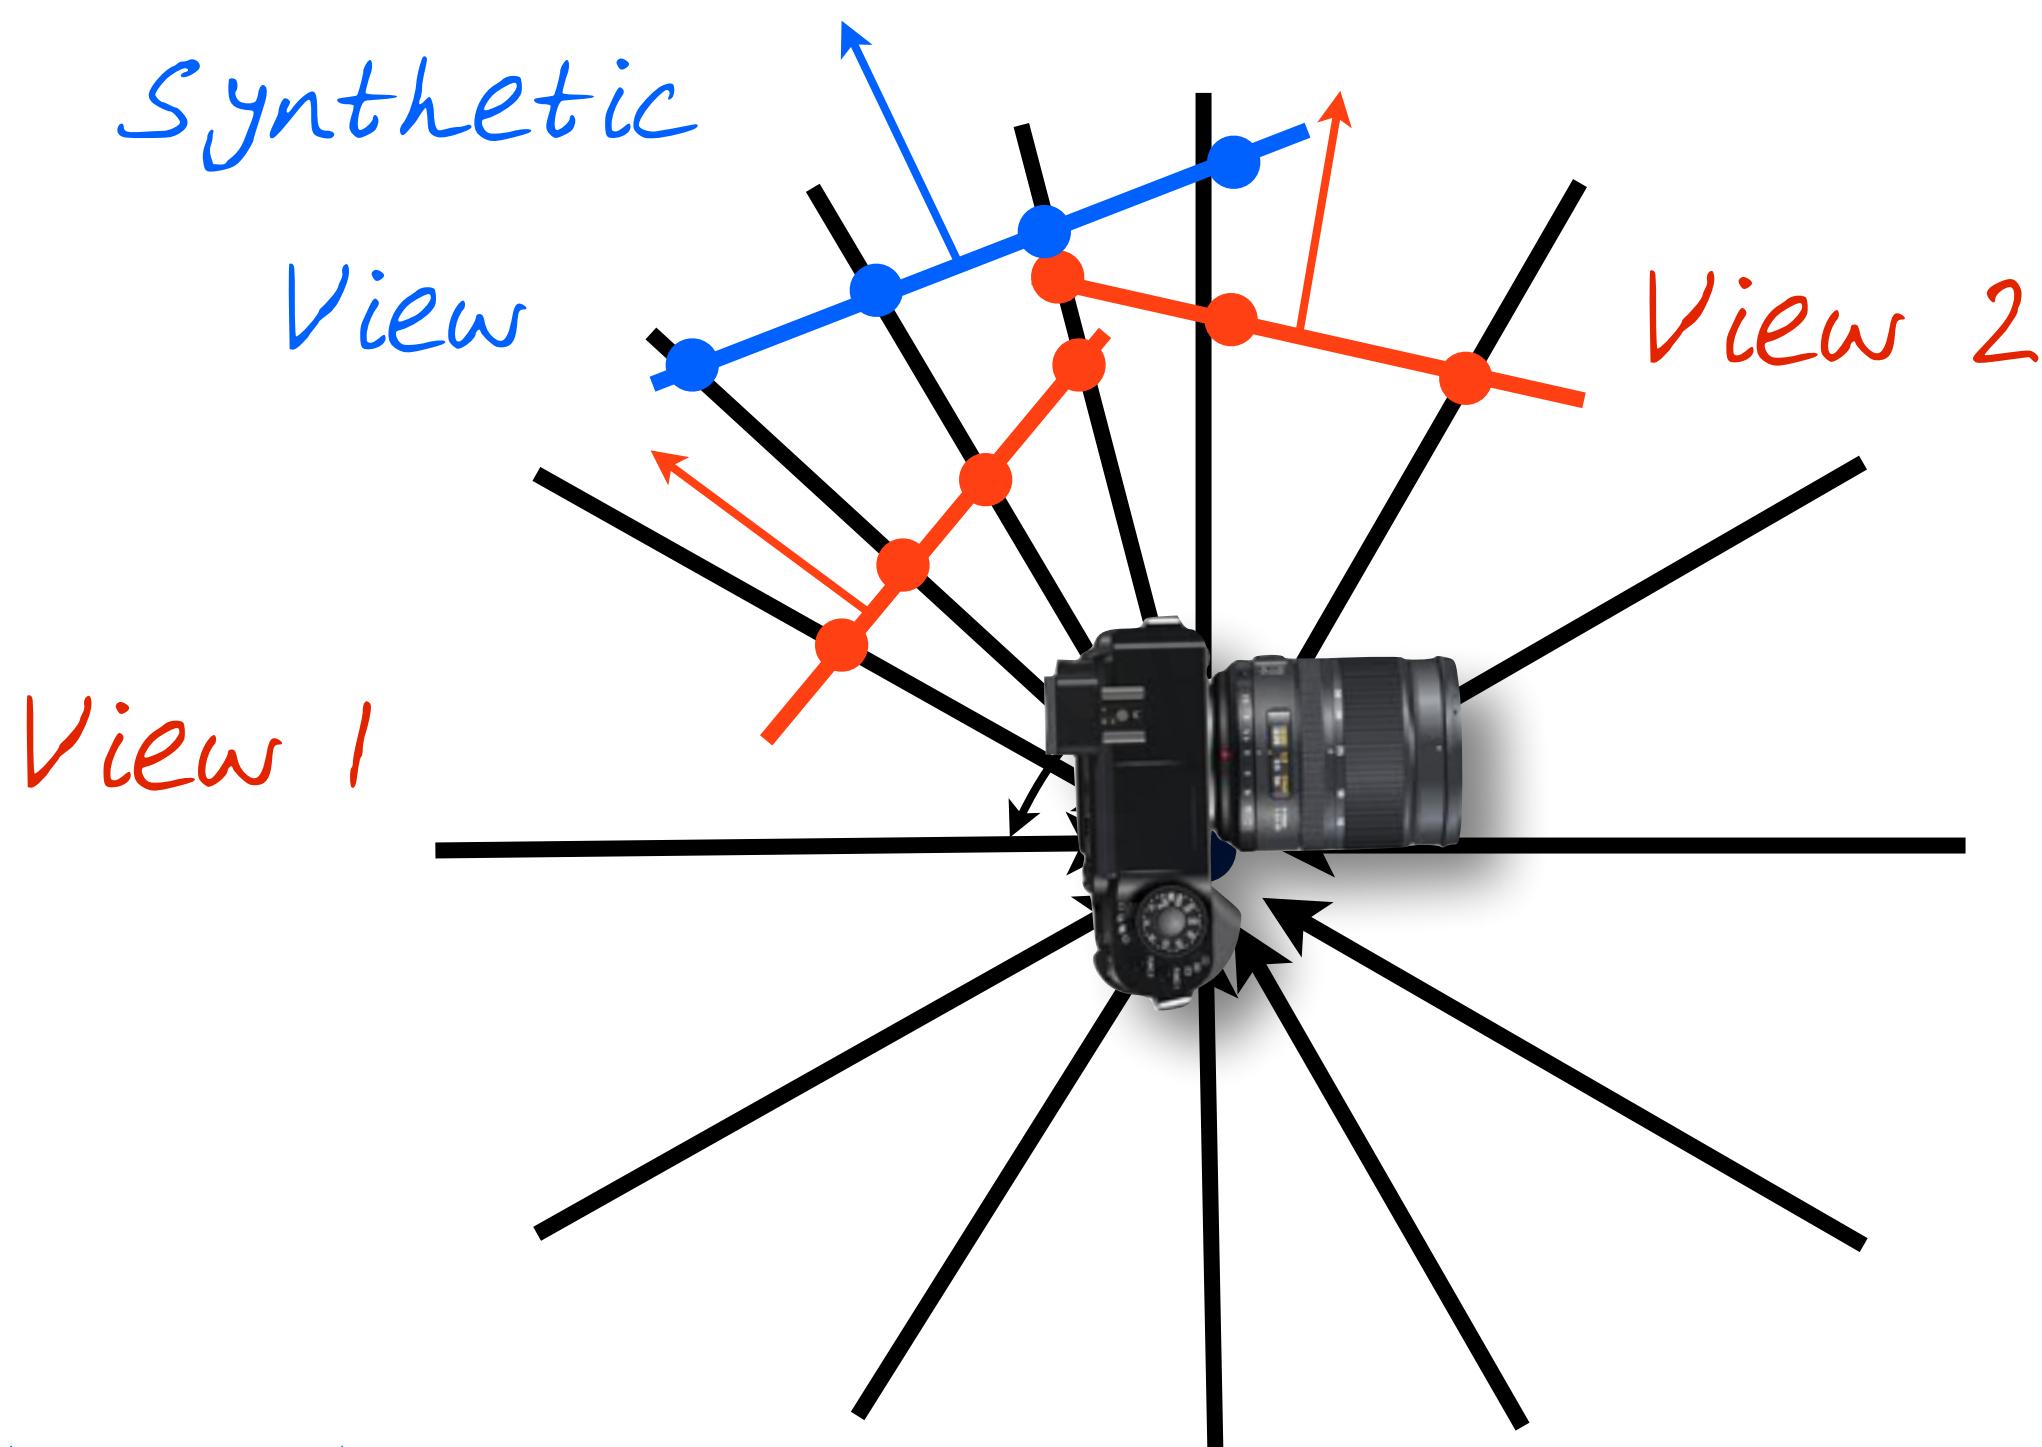


R on top

Better: Warp



A Bundle of Rays Contains all Views



Possible to generate
any synthetic
camera view as long
as it has
the same center of
projection!

Slide motivated by James Hays

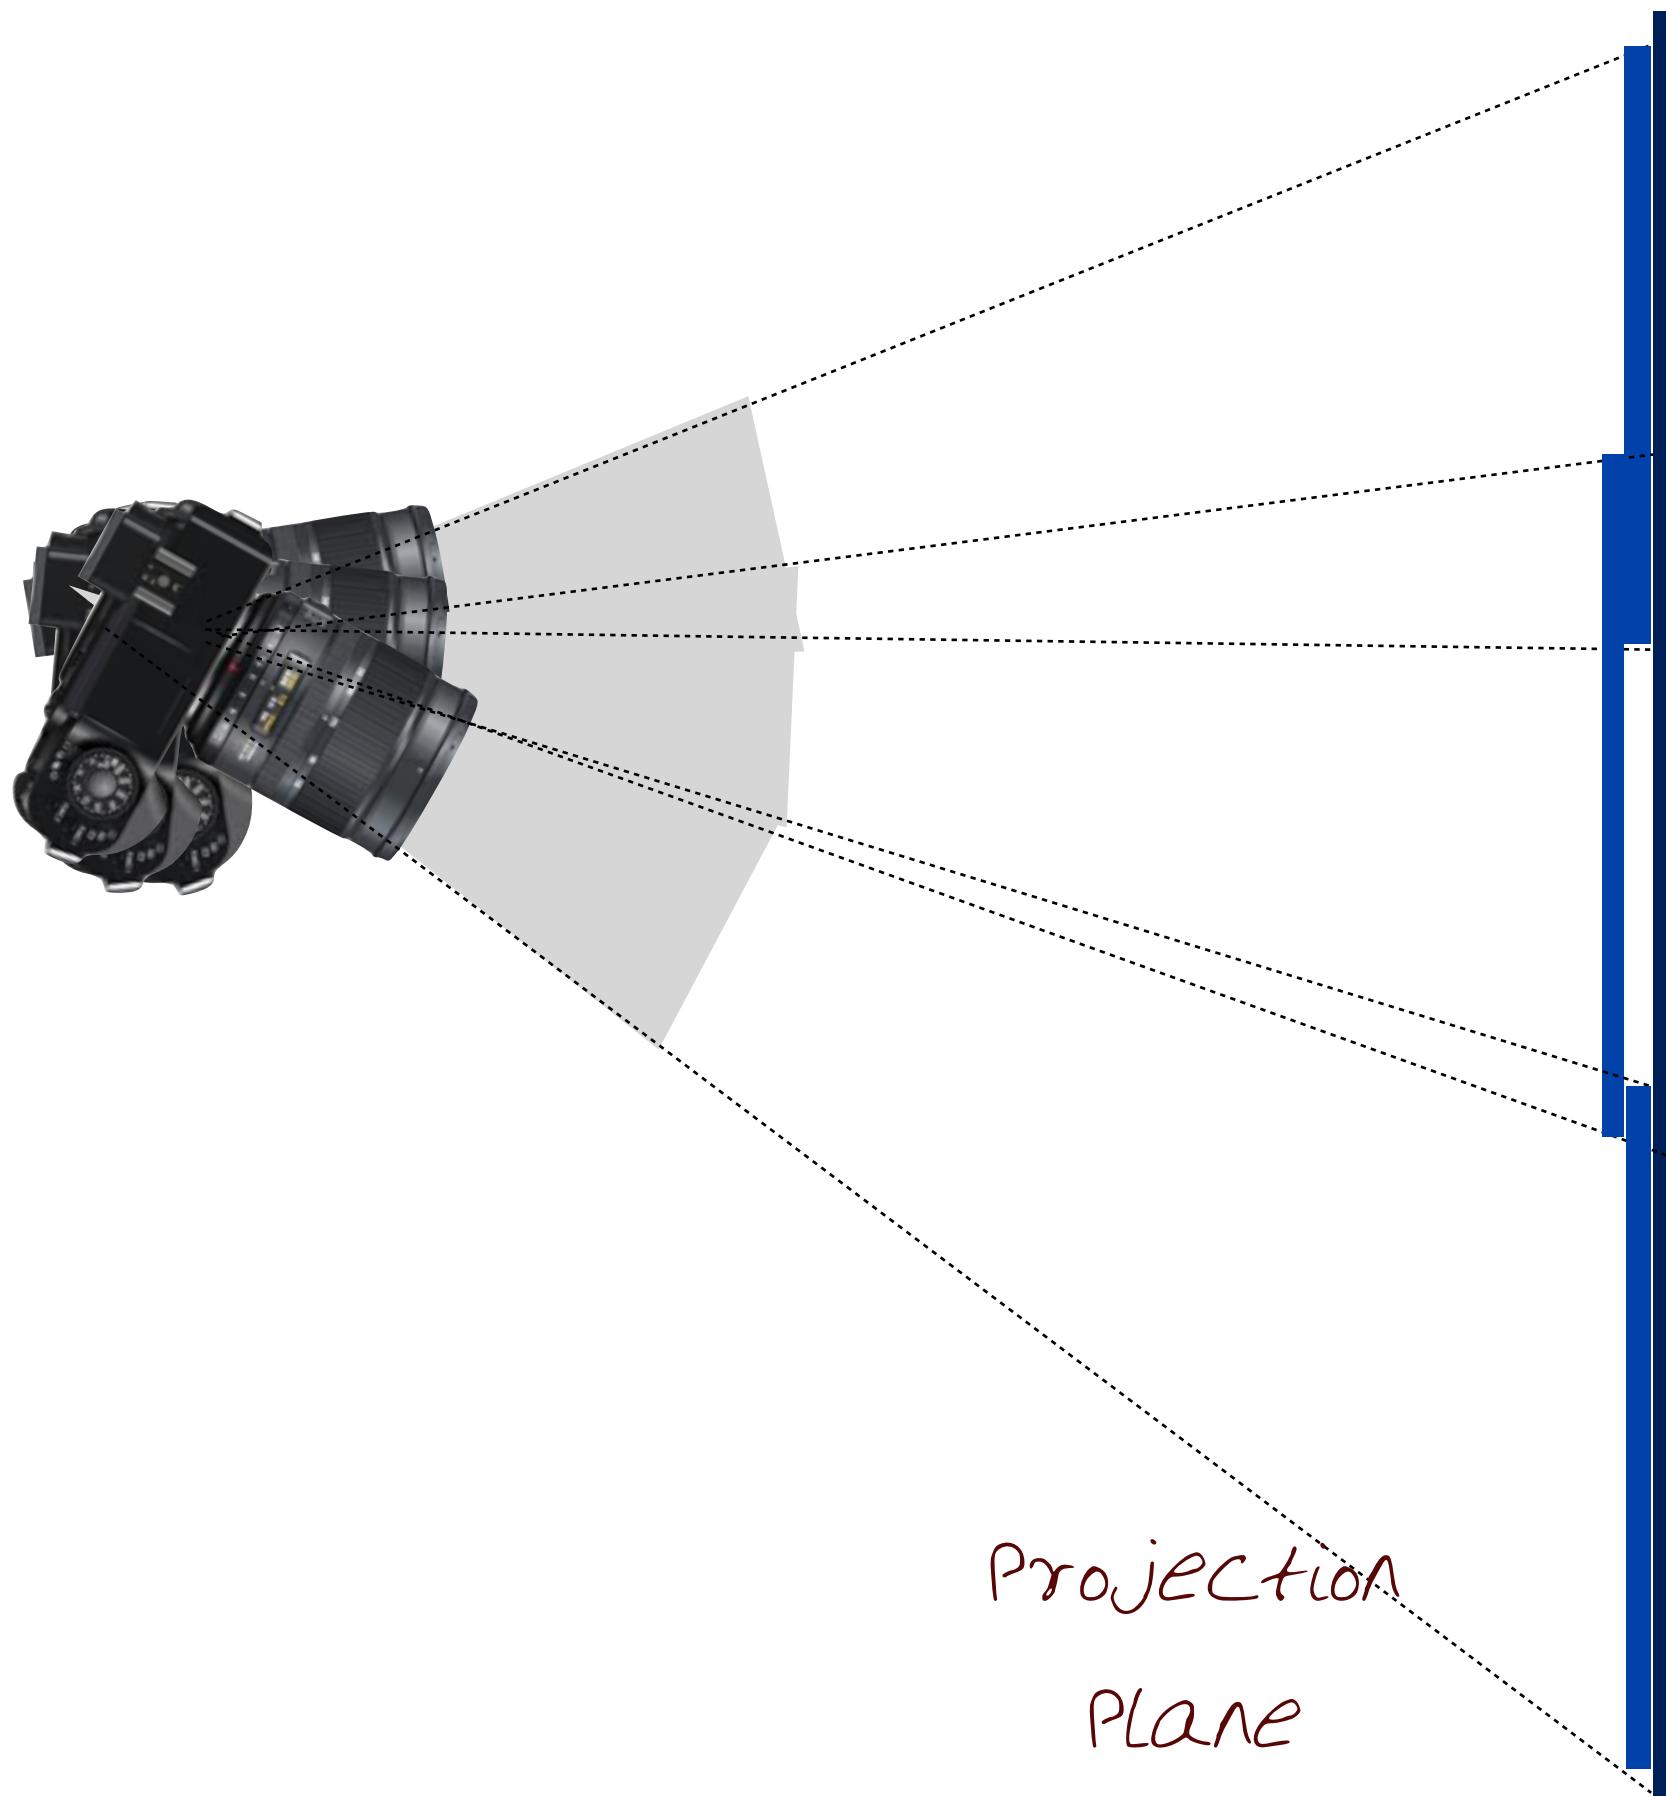


Image Re-projection to Panorama Projection Plane

- * The panorama mosaic has a natural interpretation in 3D
- * Images are reprojected onto a common plane
- * The mosaic is formed on this plane
- * Mosaic is a synthetic wide-angle camera

Image Re-Projection (I)

To relate two images from the same camera center and map a pixel from PP_1 to PP_2 :

- * Cast a ray through each pixel in PP_1
- * Draw the pixel where that ray intersects PP_2

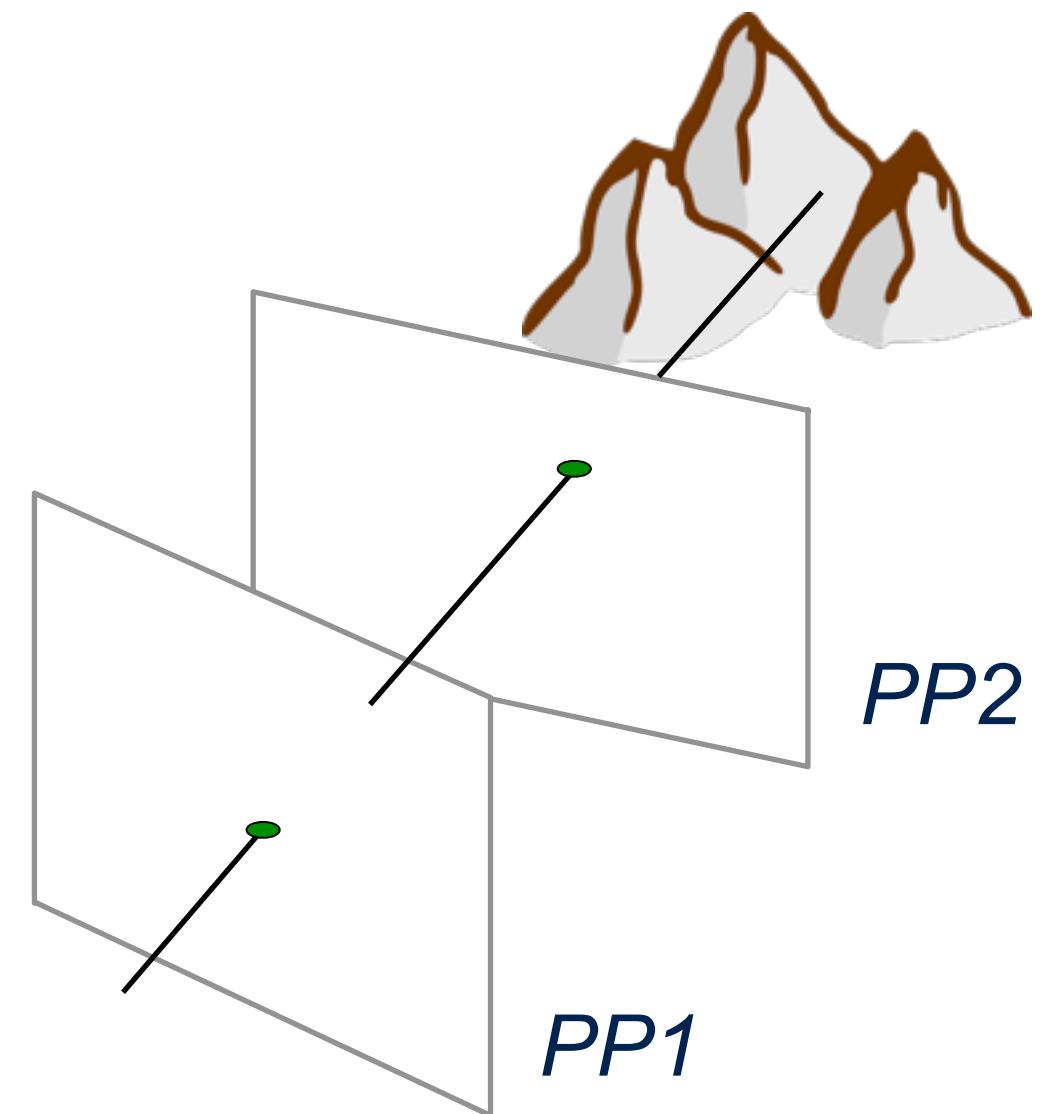
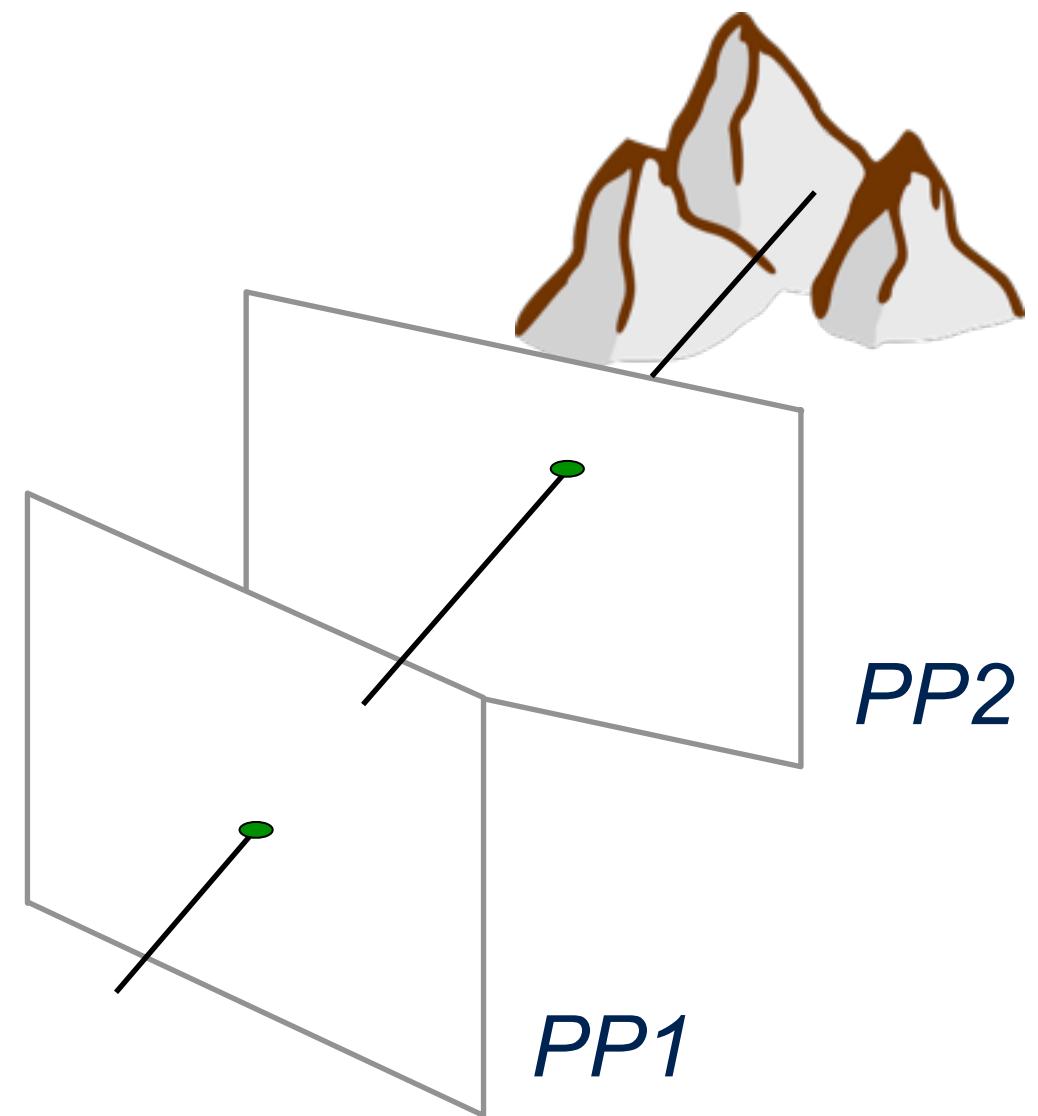


Image Re-Projection (2)

To relate two images from the same camera center and map a pixel from PP_1 to PP_2 :

- * Rather than a 3D re-projection,
- * Think of it as a 2D image warp from one image to another
- * Do not need to know the geometry of the two planes with respect to the eye?



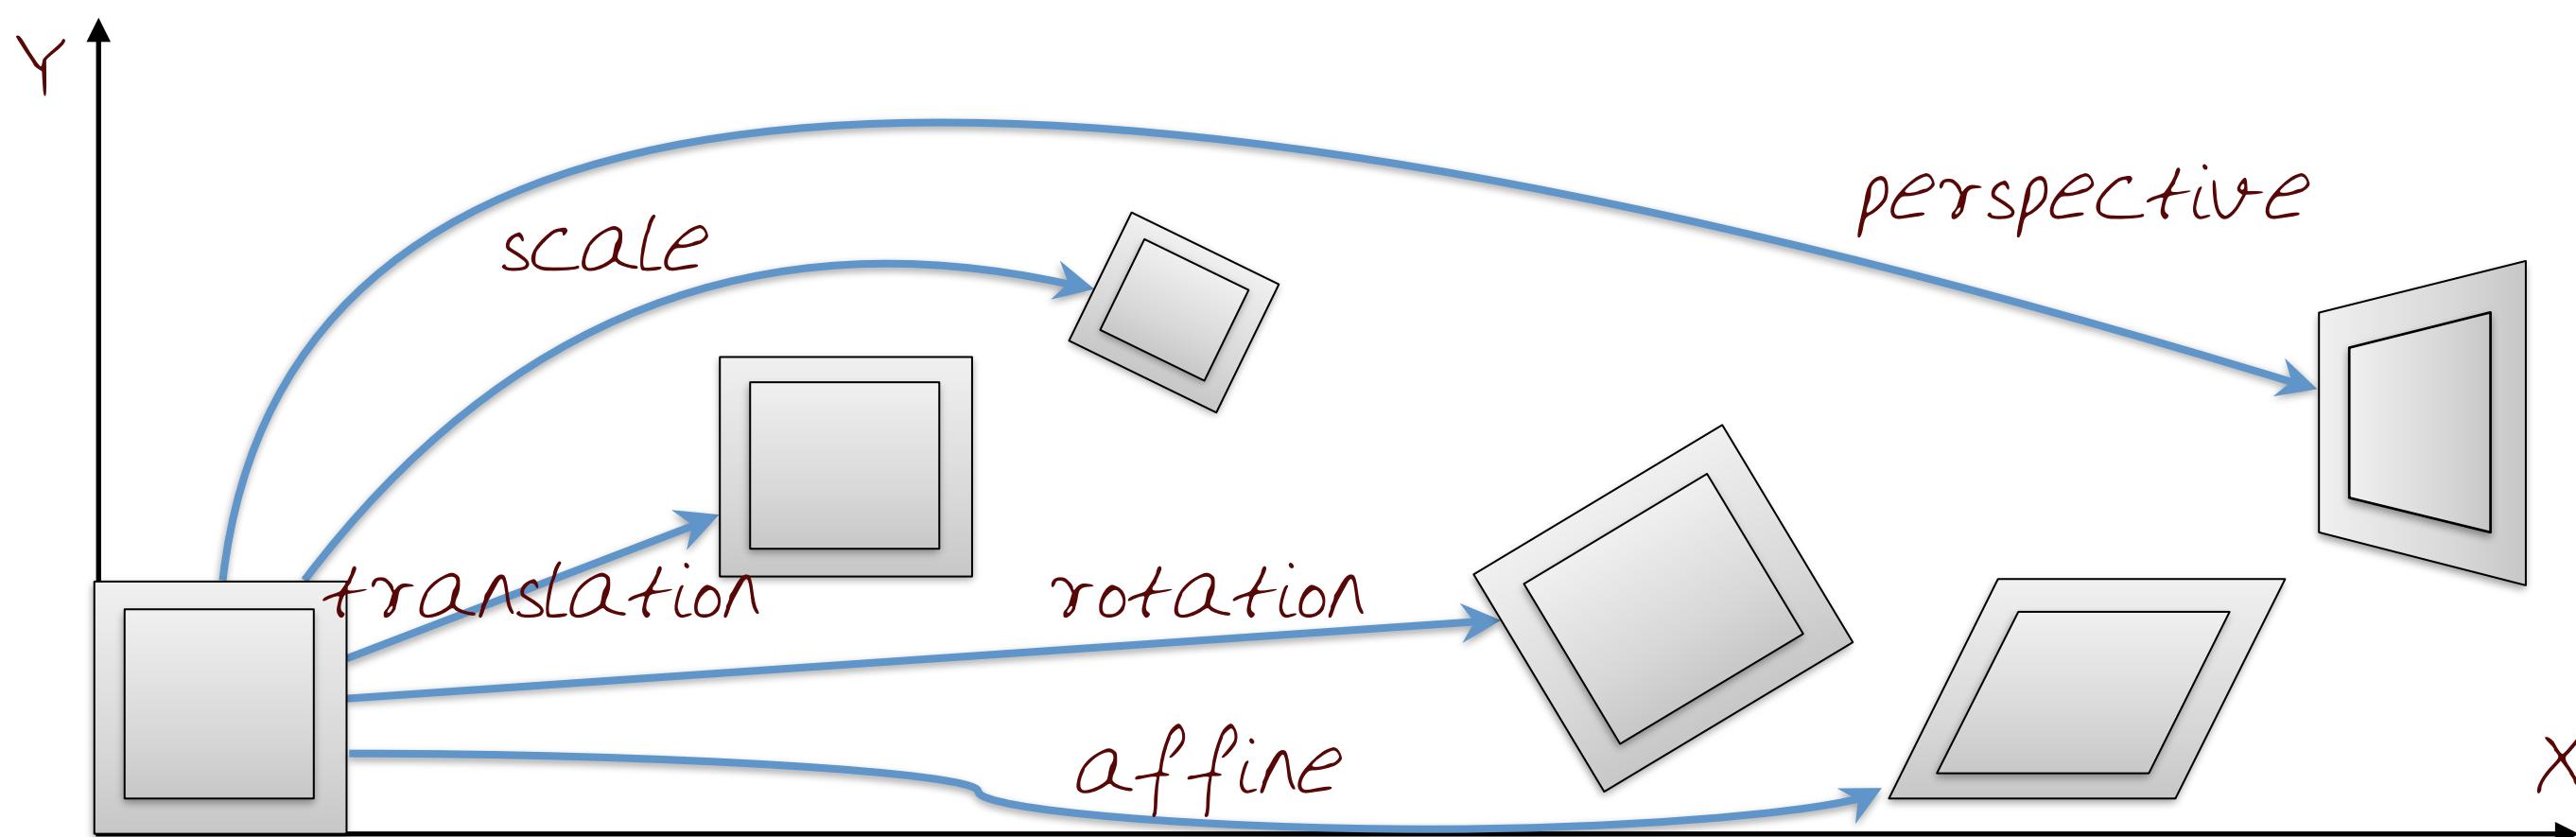
Recall: Image Warping

Which transform is the right one for warping PP_1 into PP_2 ?

E.g. translation, Euclidean, affine, projective

Translation: 2 unknowns, Euclidean: 3 unknowns

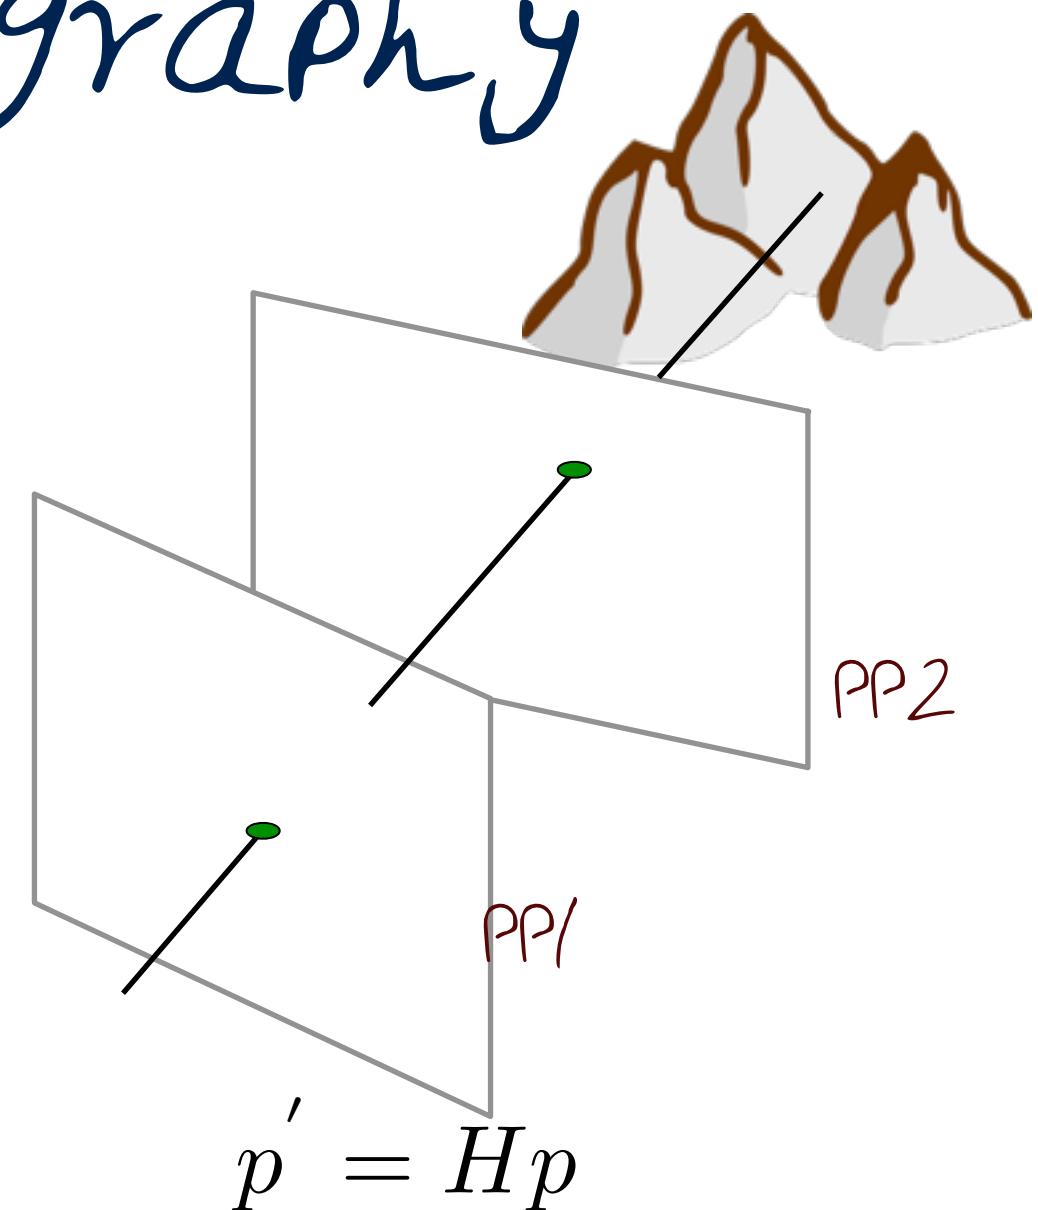
Affine: 6 unknowns, Projective: 8 unknowns



Introducing Homography

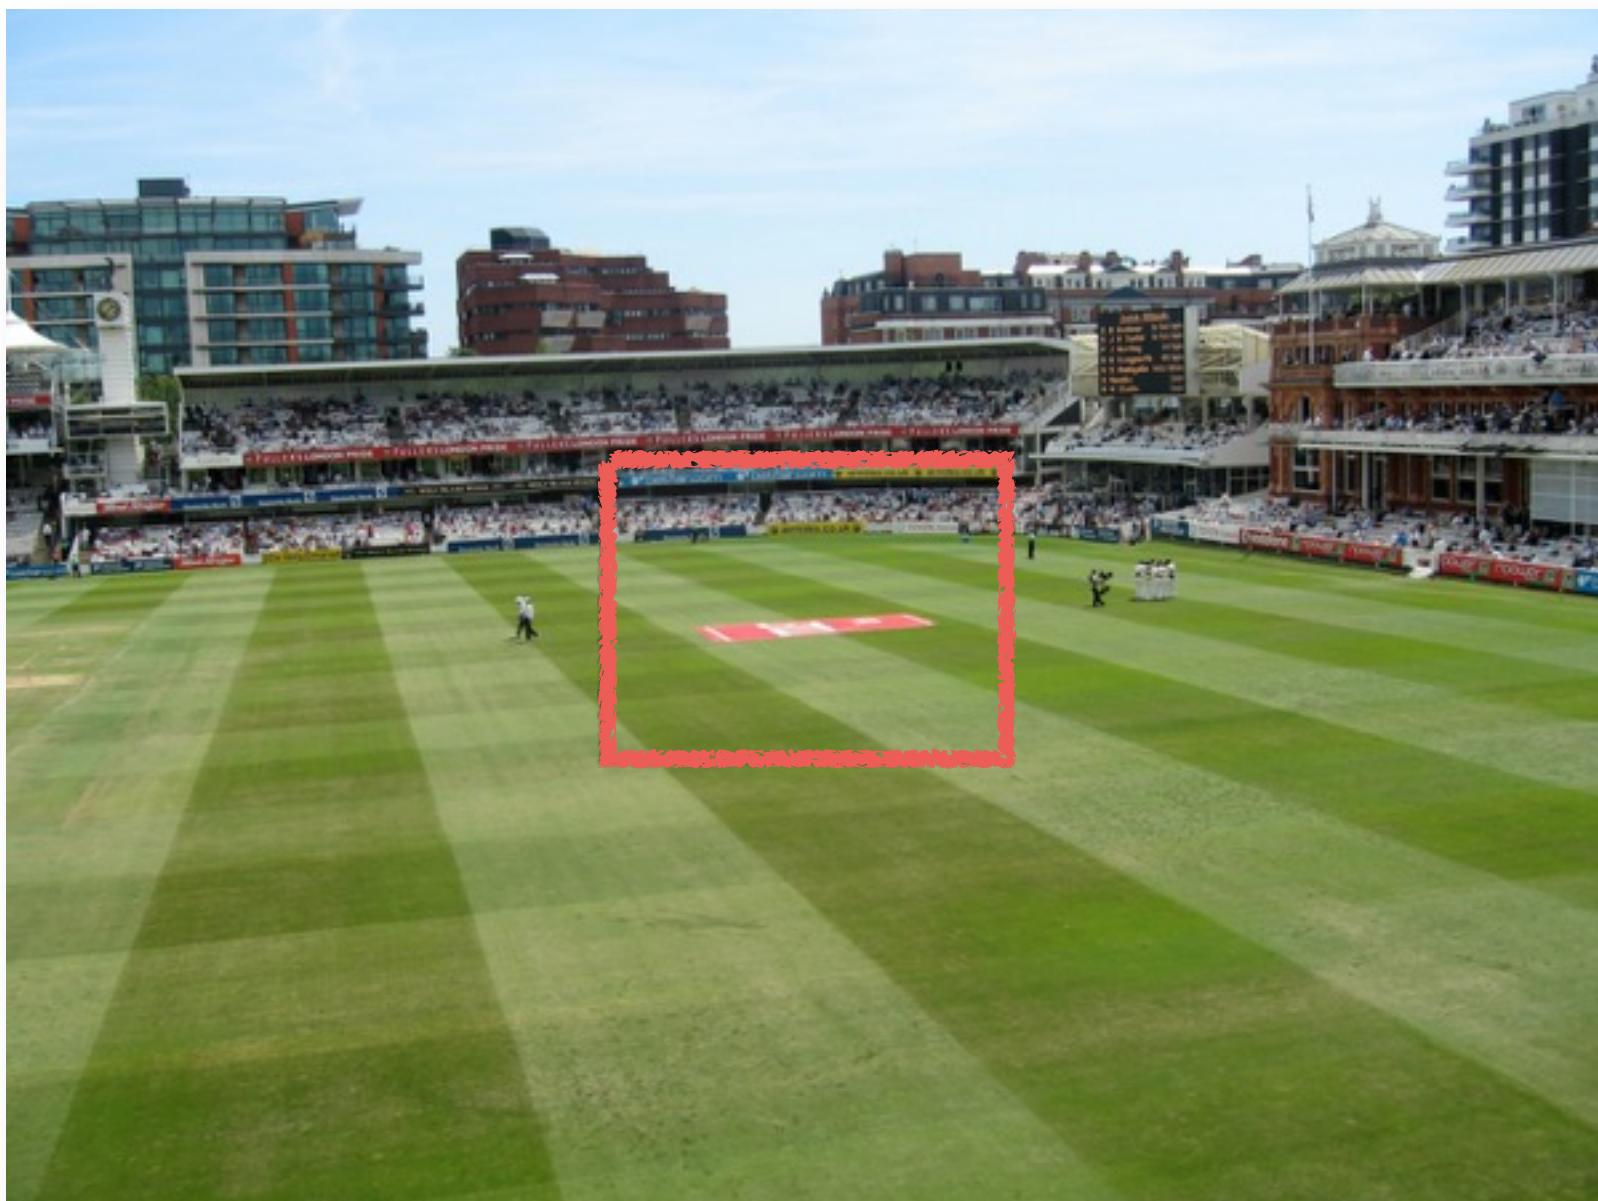
Relates two images from the same camera center

- * Rectangle should map to arbitrary quadrilateral
- * Parallel lines aren't parallel
- * Straight lines must be straight



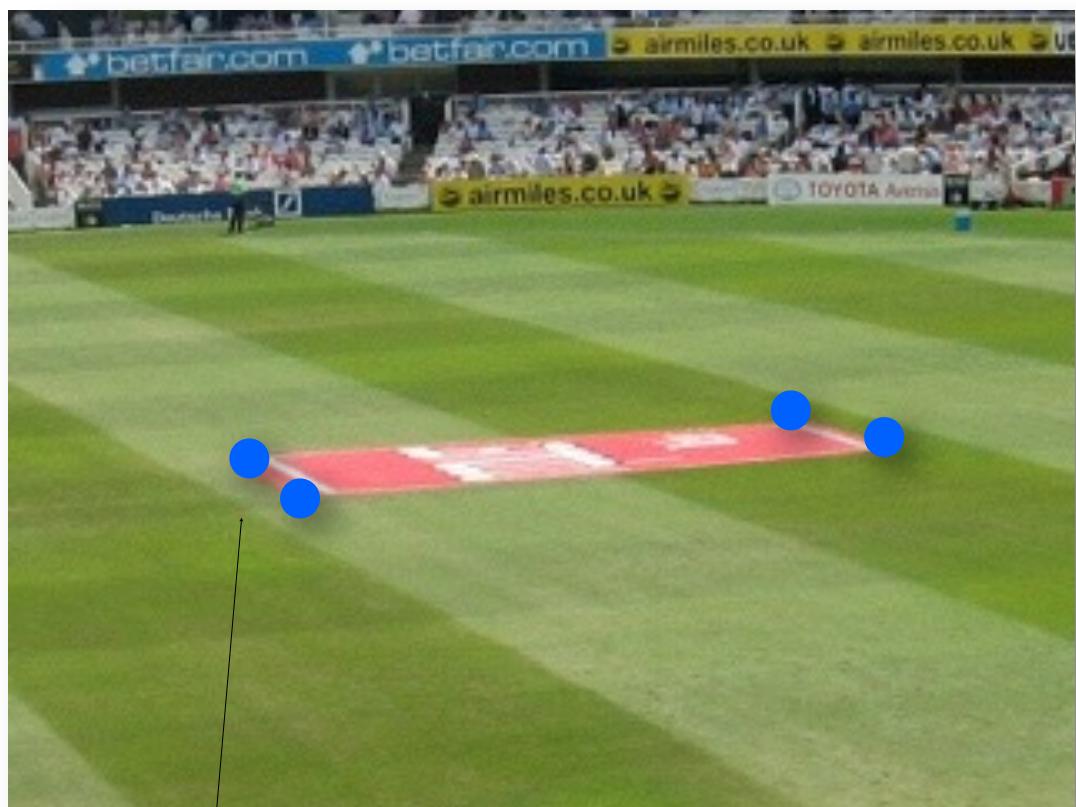
$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Computing Homography



Computing Homography

Zoomed



$$(x, y)$$

$$p_1, p_2, \dots, p_n$$

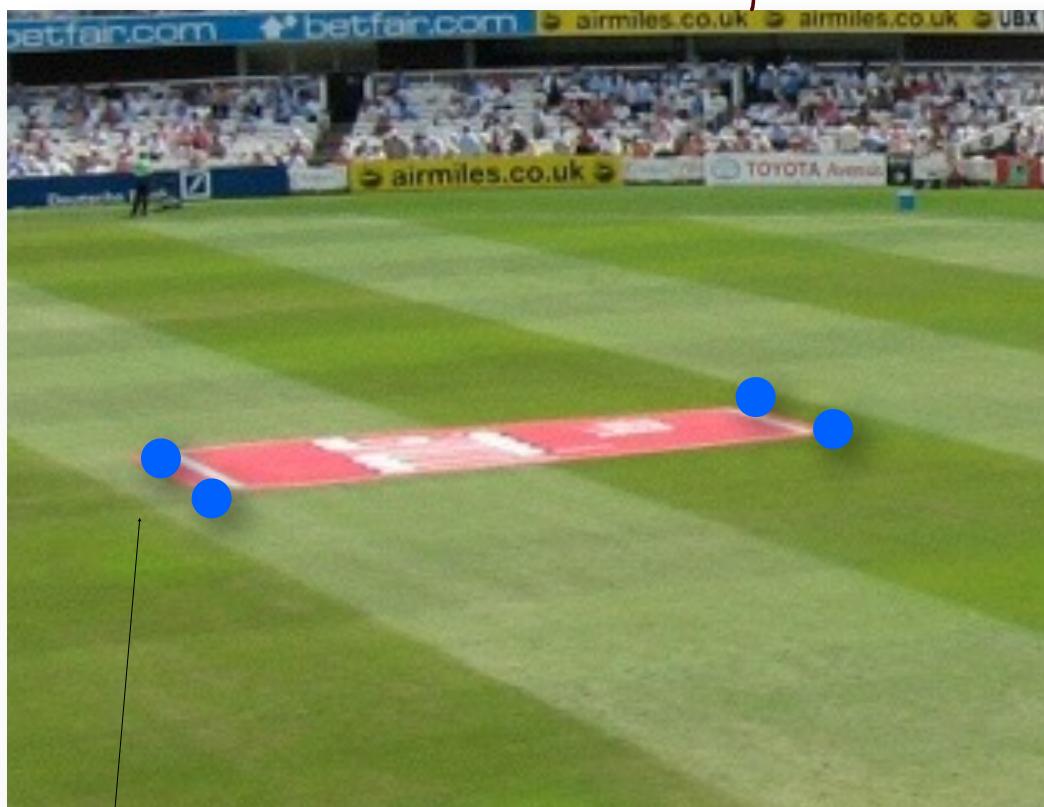
$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$(wx'/w, wy'/w) = (x', y')$$

$$p_1, p_2, \dots, p_n$$

$$p' = Hp$$

Zoomed and Left Pan



To compute the homography H , given pairs of corresponding points in two images, we need to set up an equation where the parameters of H are unknowns...

Solving for a Homography

- * Set up a system of linear equations:

$$p' = Hp$$

$$Ah = b$$

- * where vector of unknowns

$$h = [a, b, c, d, e, f, g, h]^T$$

- * Need at least 8 equations, but the more the better...

- * Solve for h . If over-constrained (i.e. more data), solve using least-squares.

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$\min \|Ah - b\|^2$$

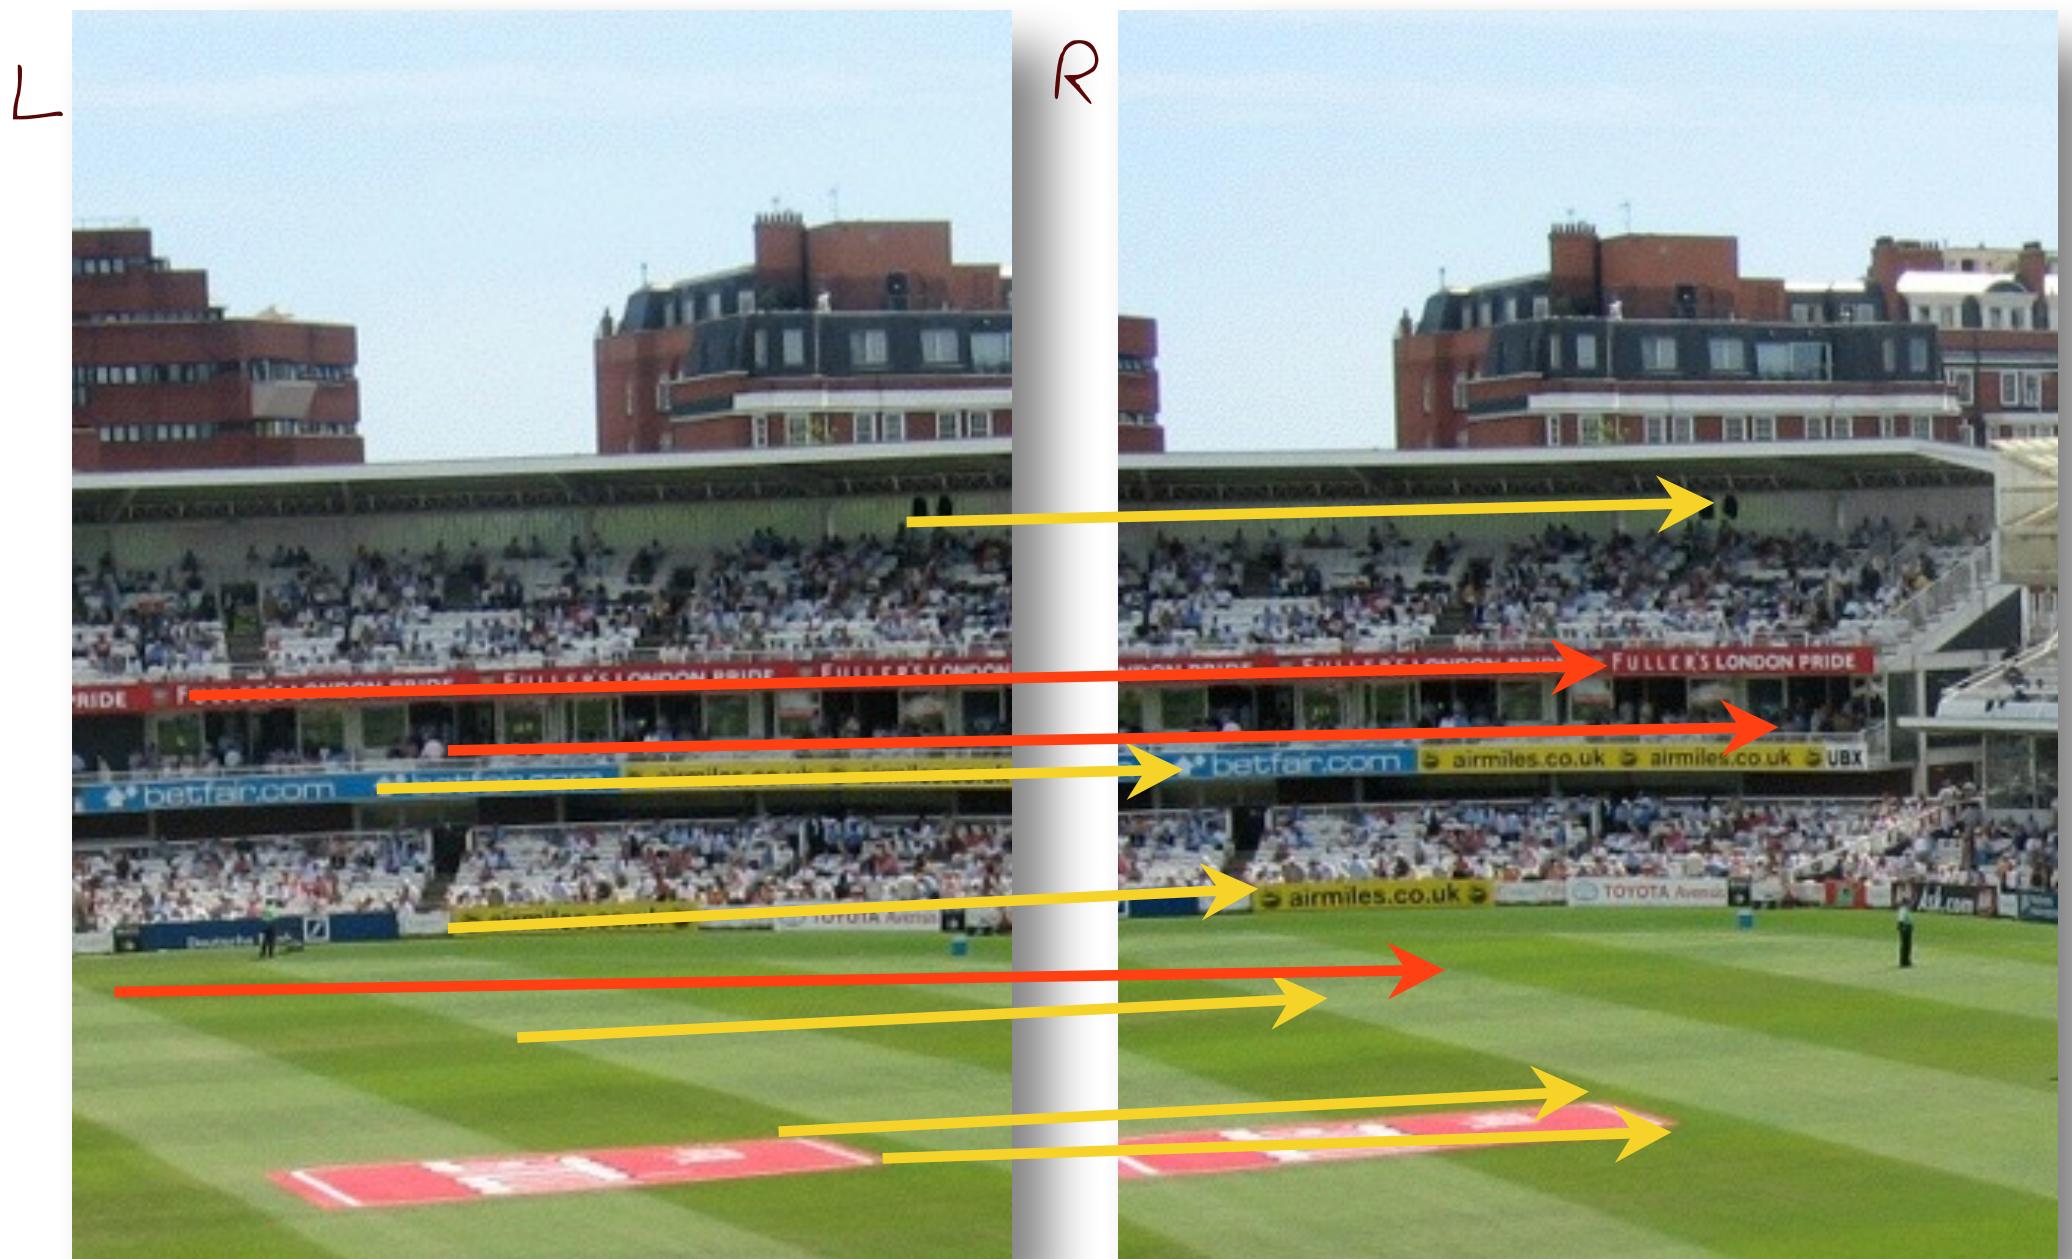
Warp into a Shared Coordinate Space



Warping and Interpolation



Dealing with BAD Matches



RANDom SAMple Consensus (RANSAC)

- * Select ONE match,
Count INLIERS
- * Find "average"
translation vector



L

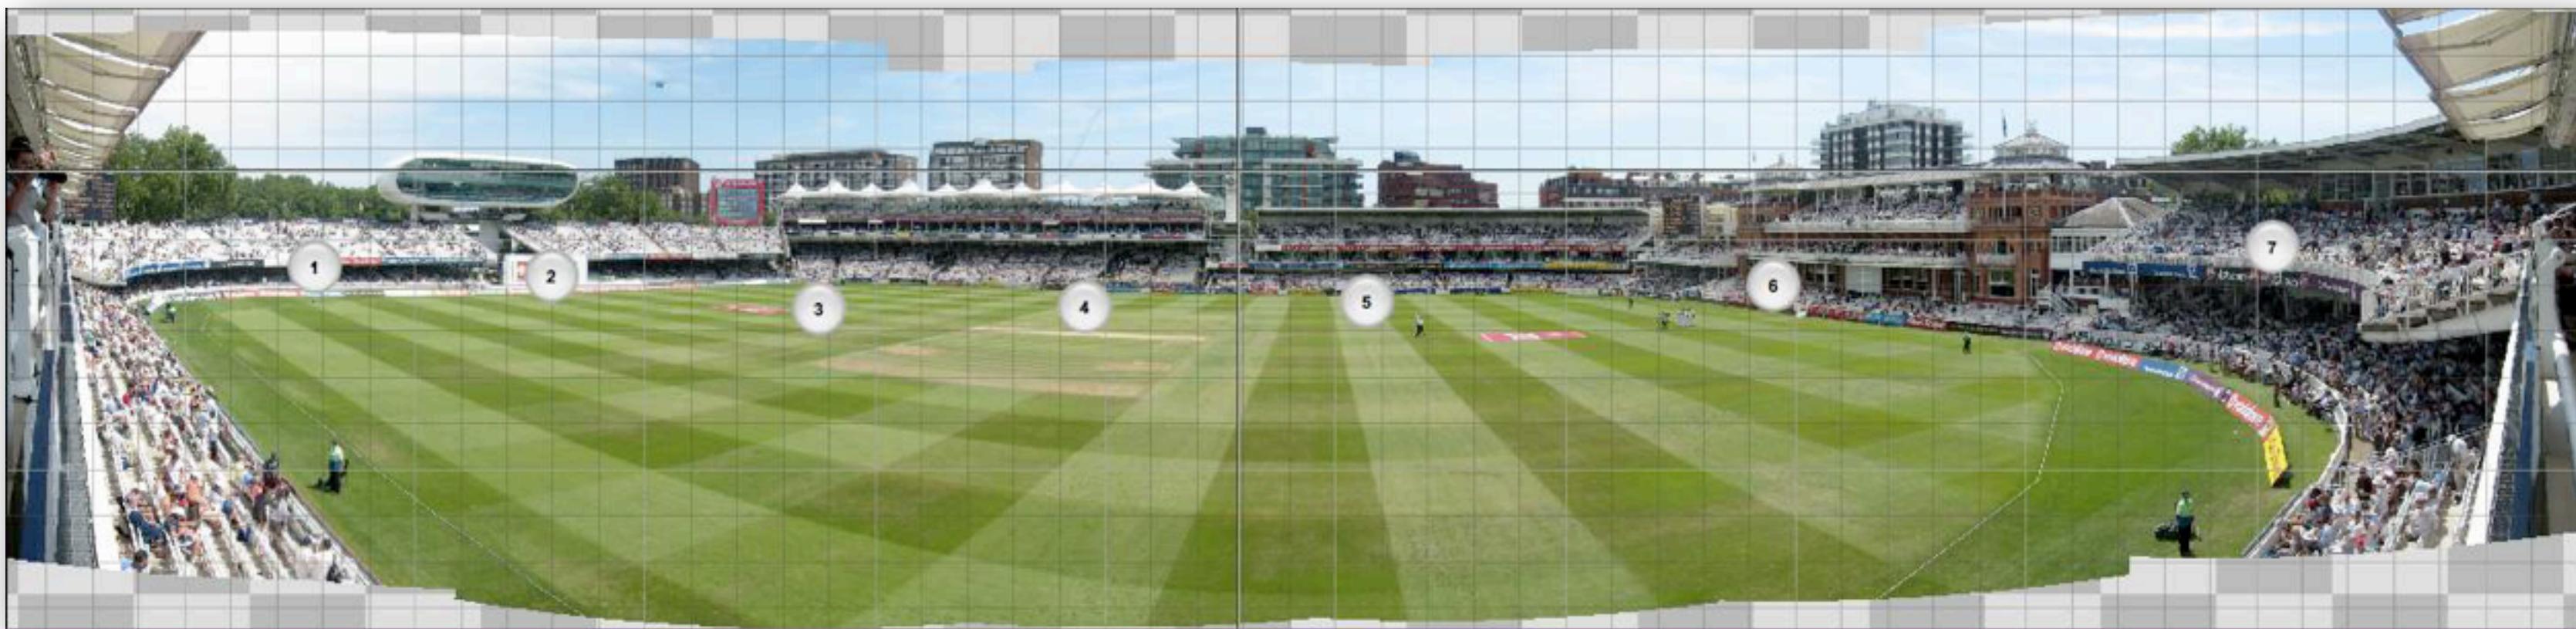
R

RANDom SAMple Consensus (RANSAC)

Loop till find a convergence/popular H :

1. Select four feature pairs (at random)
2. Compute homography H (exact)
3. Compute inliers where: $SSD(p_{in}', H p_{in}) < \varepsilon$
4. Keep largest set of inliers
5. Re-compute least-squares H estimate on all of the inliers

Key idea: Not that there are more inliers than outliers, but that the outliers are wrong in different ways.



Using kolor autopano giga™ v3



DEMO



```

import numpy as np
import cv2

# Read images
img1 = cv2.imread("einstein.png") # left image
img2 = cv2.imread("davinci.png") # right image
print "Image 1 size: {}x{}".format(img1.shape[1], img1.shape[0])
print "Image 2 size: {}x{}".format(img2.shape[1], img2.shape[0])
cv2.imshow("Image 1", img1)
cv2.imshow("Image 2", img2)

# Convert to grayscale
img1_gray = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
img2_gray = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)

# Initialize ORB detector object
orb = cv2.ORB() # or cv2.SIFT() in OpenCV 2.4.9+

# Find keypoints, compute descriptors and show them on original images (with scale and
orientation)
kp1, des1 = orb.detectAndCompute(img1, None)
kp2, des2 = orb.detectAndCompute(img2, None)
print "Image 1: {} keypoints found".format(len(kp1))
print "Image 2: {} keypoints found".format(len(kp2))
img1_kp = cv2.drawKeypoints(img1, kp1, flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
img2_kp = cv2.drawKeypoints(img2, kp2, flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
cv2.imshow("Image 1: Keypoints", img1_kp)
cv2.imshow("Image 2: Keypoints", img2_kp)

```

```

# Create two sequences of corresponding (matched) points
pts1 = []
pts2 = []
for mat in matches:
    # Get the matching keypoints for each of the images
    # Note: We could get rid of weak matches (large distance) here, but RANSAC can handle them
    pts1.append(kp1[mat.queryIdx].pt)
    pts2.append(kp2[mat.trainIdx].pt)

pts1 = np.asarray(pts1, dtype=np.float32)
pts2 = np.asarray(pts2, dtype=np.float32)

# Compute homography (since our objects are planar) using RANSAC to reject outliers
M_hom, inliers = cv2.findHomography(pts2, pts1, cv2.RANSAC) # transform img2 to img1's space
print "Computed homography (perspective transform matrix):"
print M_hom

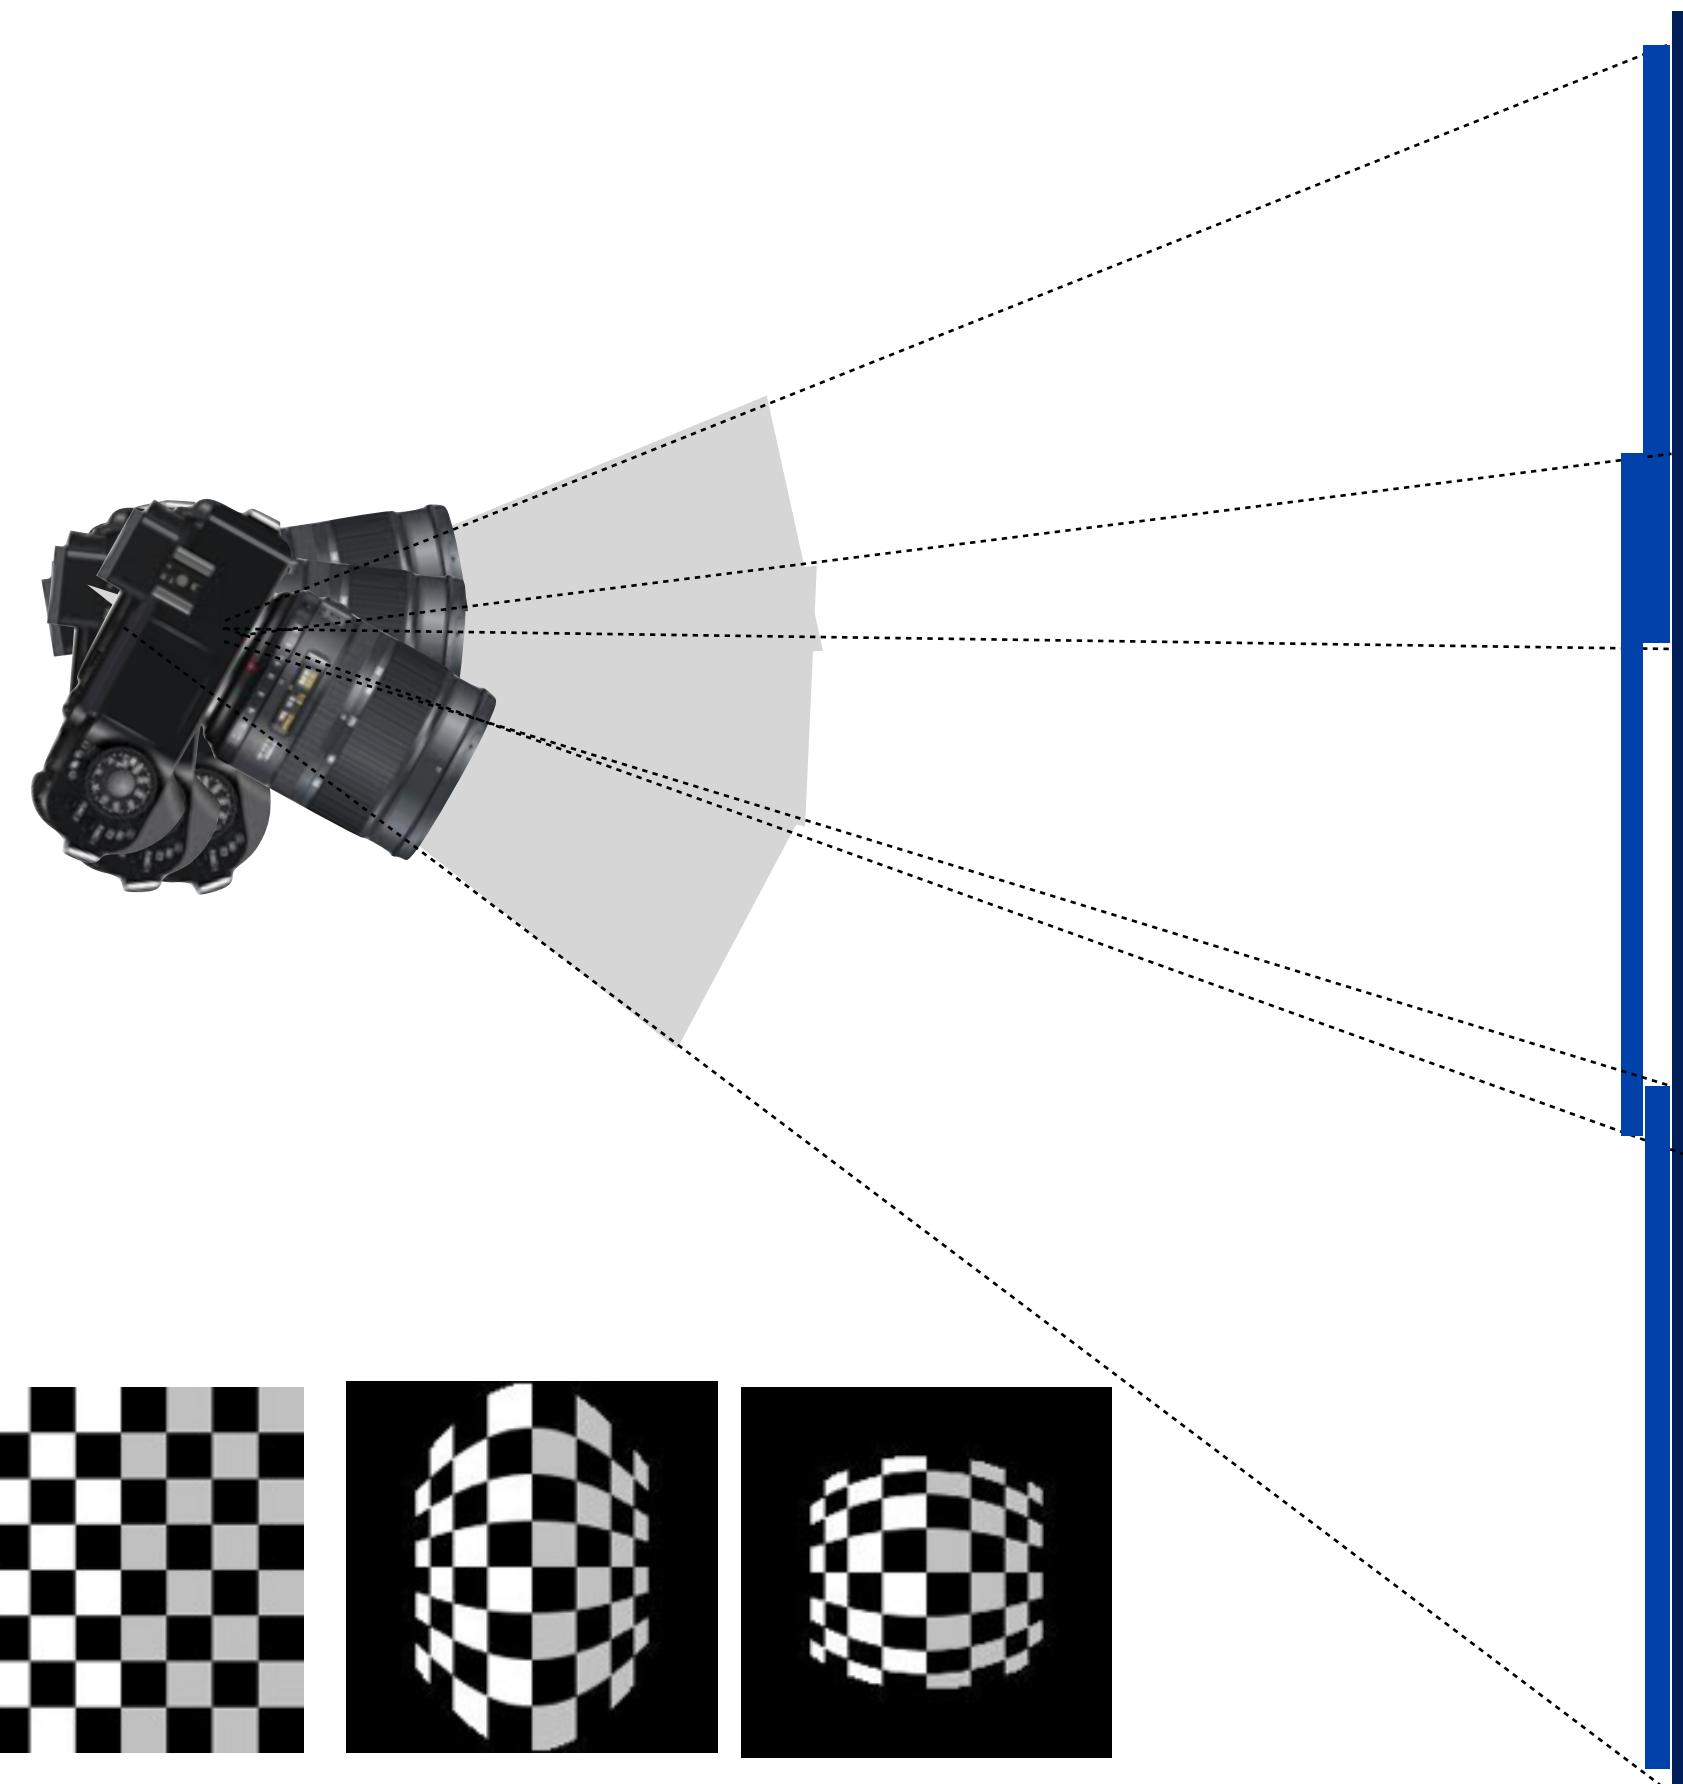
# Create panorama
pano_size = (int(M_hom[0, 2] + img2.shape[1]), max(img1.shape[0], img2.shape[0])) # combined size
img_pano = cv2.warpPerspective(img2, M_hom, pano_size) # apply transform to img2
img_pano[0:img1.shape[0], 0:img1.shape[1], :] = img1 # copy in pixels from img1

# Note: This will produce a result with visible seams; better: blend/cut/fade
cv2.imshow("Panorama", img_pano)

```

Demo





Plain Cylinder Sphere

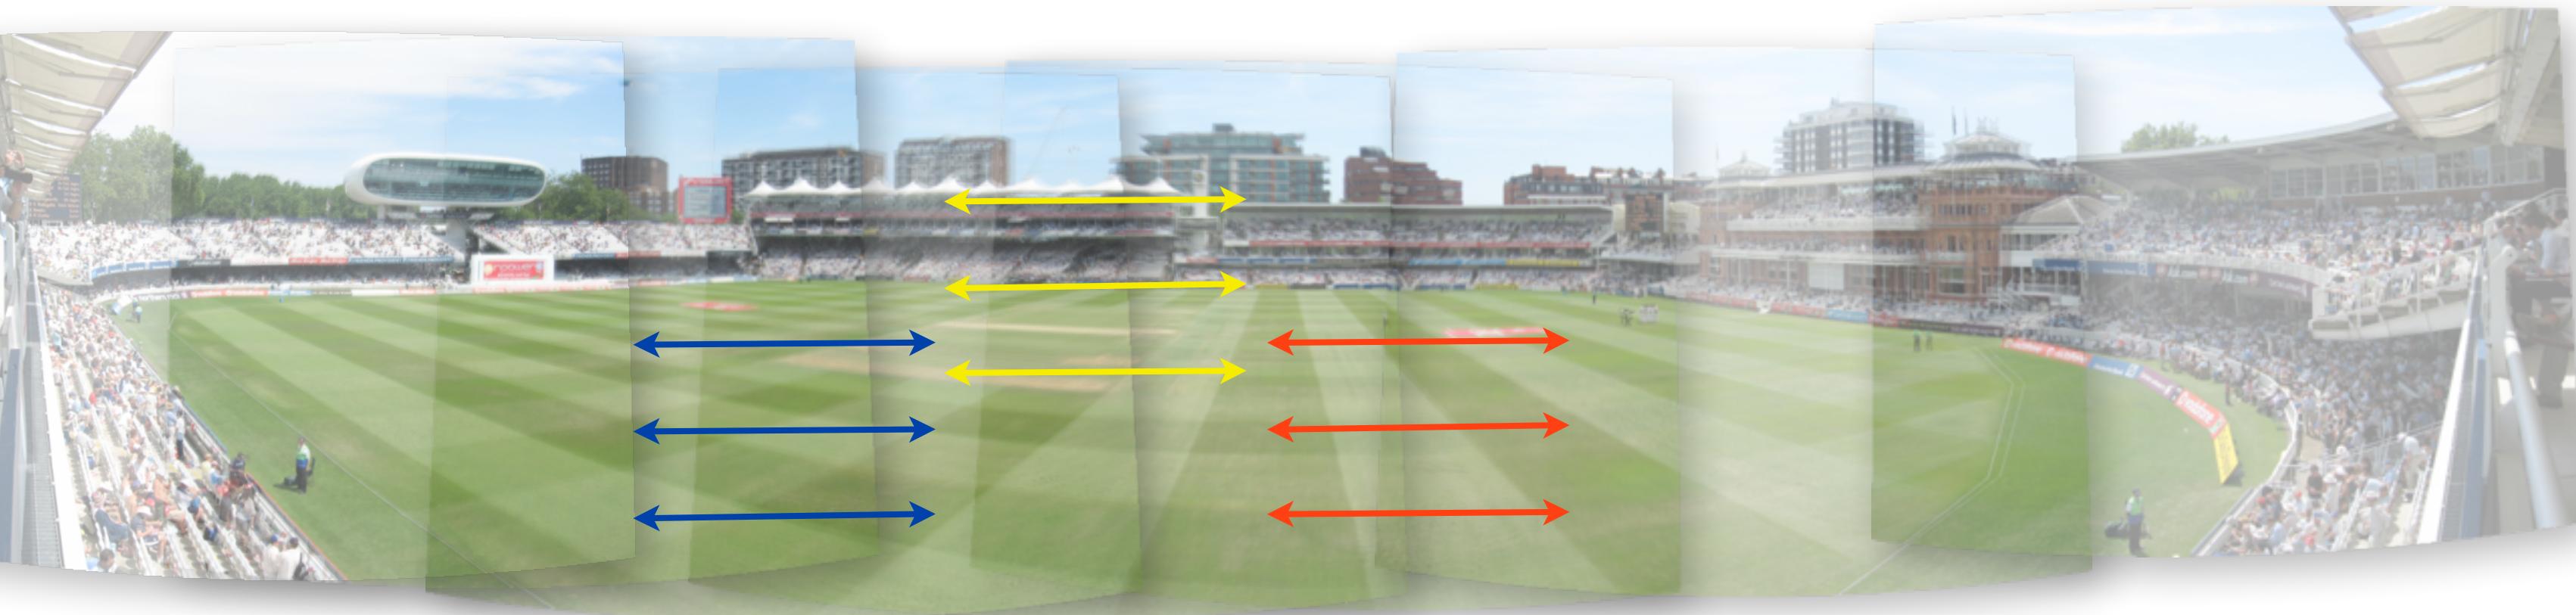
Not Just a Plane

Projection Plane

- * Cylinder
- * Sphere

Planar,
Spherical,
Cylindrical
Panoramas





“Finding Panoramas”

Using RANSAC and related matching techniques, we can find images next to each other that form a panorama. So we don't have to take pictures in a sequence.

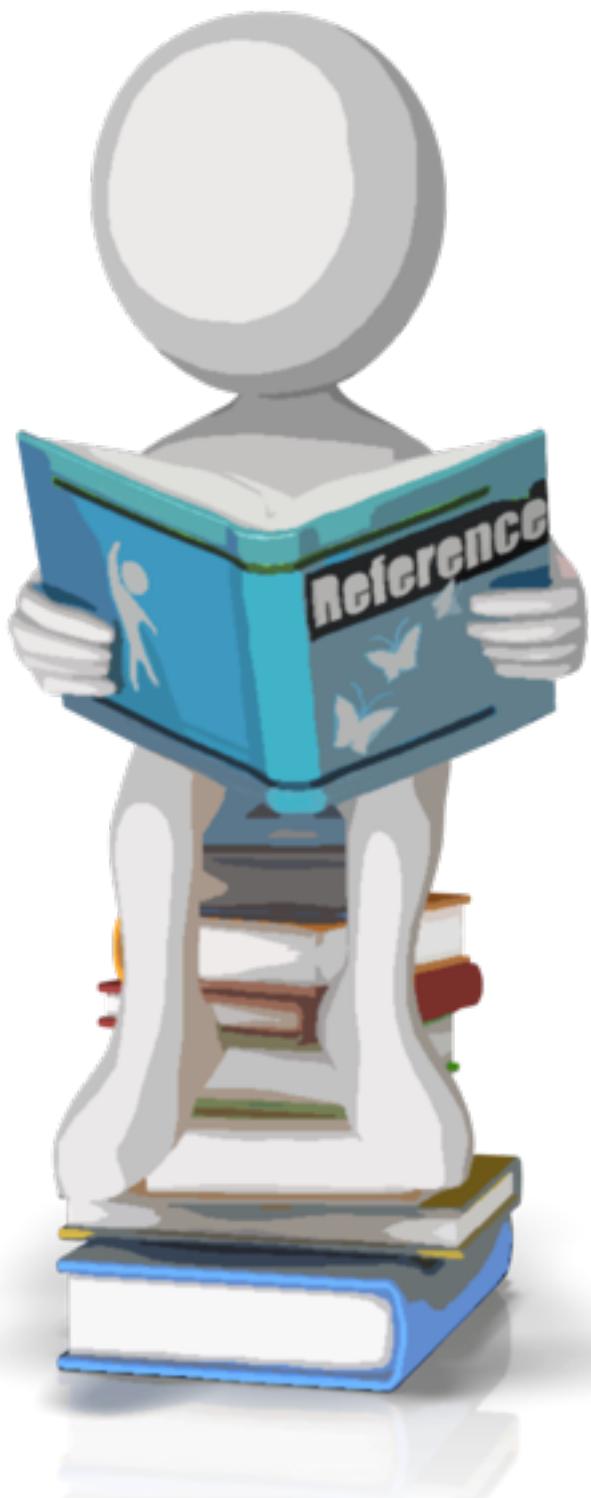
Brown and Lowe (2003).

Summary

- * Five (5) steps to generate a Panorama
- * Image Re-projection for Panoramas
- * Homography and how it is computed from a pair of images
- * RANSAC as an approach to deal with Bad matches across images
- * Additional considerations for Panoramas



Further Reading



- * Brown and Lowe (2003). "Recognising Panoramas." International Conference on Computer Vision (ICCV2003) ([pdf](#) | [bib](#) | [ppt](#))
- * Microsoft Research Image Composite Editor (ICE)
- * Panorama Tools Graphical User Interface (PTGui)
- * [Hugin Panorama Photo Stitcher](#)

Computational Photography

- * Study the basics of computation and its impact on the entire workflow of photography, from capturing, manipulating and collaborating on, and sharing photographs.



© 2015 Irfan Essa, Georgia Tech, All Rights Reserved